

Practical No 6

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
df = pd.read_csv("xAPI-Edu-Data.csv")
print("\n==== First 5 rows of dataset ===")
print(df.head())
# Step 1: Scan Missing Values & Inconsistencies
print("\n==== Missing Values ===")
print(df.isnull().sum())
# Fill numeric missing with median
for col in df.select_dtypes(include=['int64','float64']).columns:
    df[col].fillna(df[col].median(), inplace=True)
# Fill categorical missing with mode
for col in df.select_dtypes(include='object').columns:
    df[col].fillna(df[col].mode()[0], inplace=True)
    df[col] = df[col].str.strip().str.lower() # fix inconsistencies
print("\n==== After Missing Value Treatment ===")
print(df.isnull().sum())
# Step 2: Outlier Detection & Treatment (IQR Method)
numeric_cols = ['raisedhands', 'VisITedResources', 'AnnouncementsView', 'Discussion']
for col in numeric_cols:
    Q1 = df[col].quantile(0.25)
    Q3 = df[col].quantile(0.75)
    IQR = Q3 - Q1
    lower = Q1 - 1.5 * IQR
    upper = Q3 + 1.5 * IQR
```

```

# Cap outliers (winsorization)
df[col] = np.where(df[col] > upper, upper,
                   np.where(df[col] < lower, lower, df[col]))

print("\n==== Outliers handled using IQR capping ===")

# Step 3: Data Transformation

print("\n==== Distribution before transformation (raisedhands) ===")
sns.histplot(df['raisedhands'], kde=True)

plt.show()

df['raisedhands_log'] = np.log1p(df['raisedhands'])

print("\n==== Distribution after log transformation (raisedhands) ===")
sns.histplot(df['raisedhands_log'], kde=True)

plt.show()

print("\n==== Final Dataset Shape ===")
print(df.shape)

```

Output:-

```

==== First 5 rows of dataset ===
   gender NationalITY PlaceofBirth      StageID GradeID SectionID Topic \
0       M           KW      KuwaIT  lowerlevel    G-04          A     IT
1       M           KW      KuwaIT  lowerlevel    G-04          A     IT
2       M           KW      KuwaIT  lowerlevel    G-04          A     IT
3       M           KW      KuwaIT  lowerlevel    G-04          A     IT
4       M           KW      KuwaIT  lowerlevel    G-04          A     IT

   Semester Relation  raisedhands  VisITEDResources AnnouncementsView \
0         F    Father        15              16                  2
1         F    Father        20              20                  3
2         F    Father        10                 7                  0
3         F    Father        30              25                  5
4         F    Father        40              50                 12

   Discussion ParentAnsweringSurvey ParentschoolSatisfaction \
0            20                      Yes                    Good
1            25                      Yes                    Good
2            30                      No                     Bad
3            35                      No                     Bad
4            50                      No                     Bad

   StudentAbsenceDays Class
0                Under-7      M

```

```
1      Under-7      M
2      Above-7     L
3      Above-7     L
4      Above-7     M
```

==== Missing Values ===

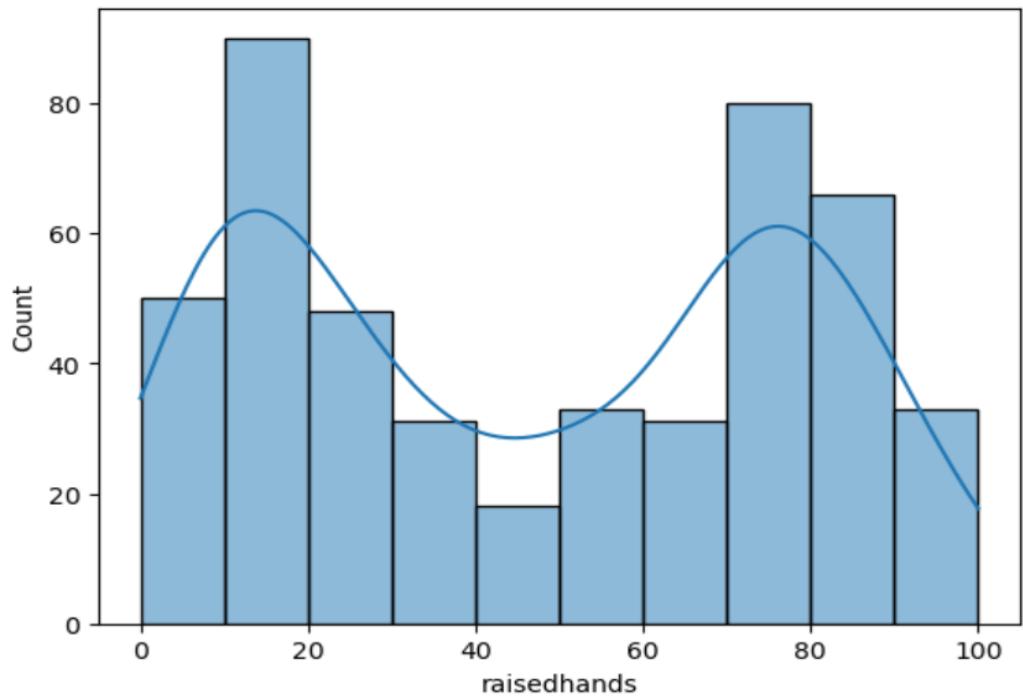
```
gender                  0
NationalITY              0
PlaceofBirth              0
StageID                  0
GradeID                  0
SectionID                 0
Topic                     0
Semester                  0
Relation                  0
raisedhands                0
VisITEDResources          0
AnnouncementsView         0
Discussion                  0
ParentAnsweringSurvey       0
ParentschoolSatisfaction    0
StudentAbsenceDays         0
Class                      0
dtype: int64
```

==== After Missing Value Treatment ===

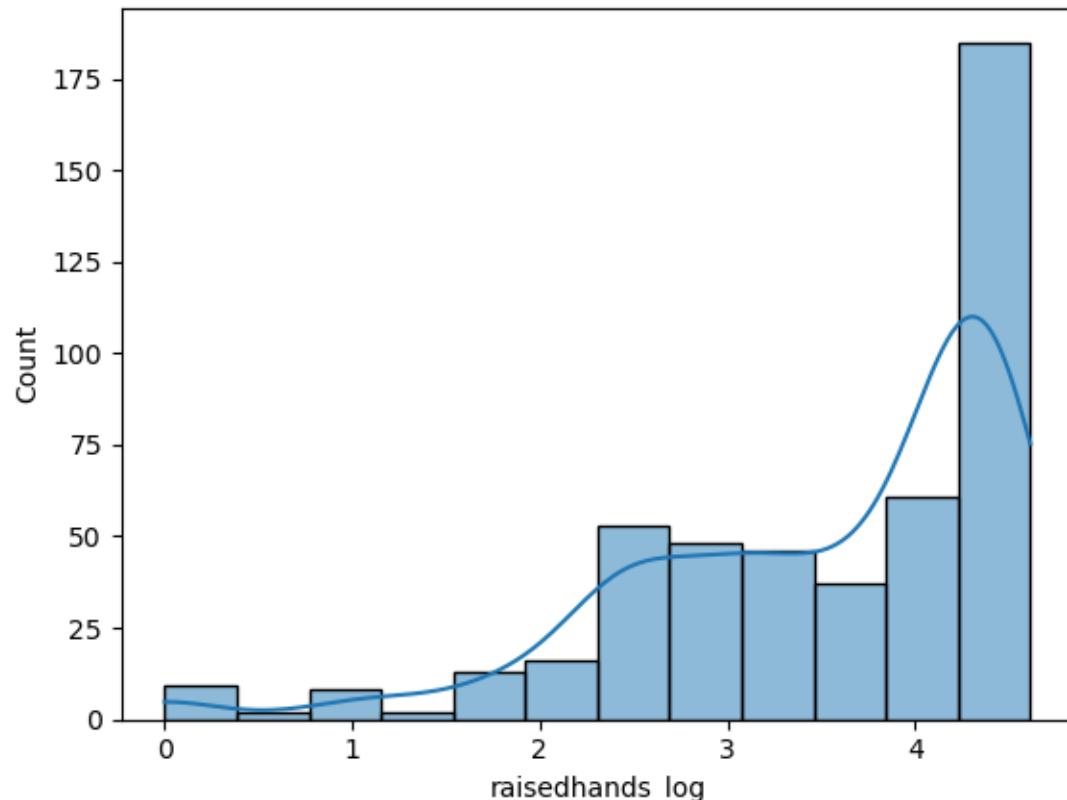
```
gender                  0
NationalITY              0
PlaceofBirth              0
StageID                  0
GradeID                  0
SectionID                 0
Topic                     0
Semester                  0
Relation                  0
raisedhands                0
VisITEDResources          0
AnnouncementsView         0
Discussion                  0
ParentAnsweringSurvey       0
ParentschoolSatisfaction    0
StudentAbsenceDays         0
Class                      0
dtype: int64
```

==== Outliers handled using IQR capping ===

==== Distribution before transformation (raisedhands) ===



==== Distribution after log transformation (raisedhands) ====



==== Final Dataset Shape ====

(480, 18)

Practical No 7

```
import pandas as pd

df = pd.read_csv("Bigmart_Sales.csv")

print("== First 5 rows ==")

print(df.head())

# 2. Pandas Indexing

print("\n== Column Access ==")

print(df['Item_Identifier'].head()) # single column

print("\n== Row Access with iloc (first 5 rows, first 3 cols) ==")

print(df.iloc[0:5, 0:3])

print("\n== Row Access with loc (rows where Outlet_Type = 'Supermarket Type1') ==")

print(df.loc[df['Outlet_Type'] == 'Supermarket Type1'].head())

# 3. Handling Null Values

print("\n== Null Values Count ==")

print(df.isnull().sum())

# Fill numeric null values with mean

for col in df.select_dtypes(include=['int64','float64']).columns:

    df[col].fillna(df[col].mean(), inplace=True)

# Fill categorical null values with mode

for col in df.select_dtypes(include='object').columns:

    df[col].fillna(df[col].mode()[0], inplace=True)

print("\n== After Filling Null Values ==")

print(df.isnull().sum())

# Drop rows if still any null values remain

df.dropna(inplace=True)

# 4. Merging, Joining, Concatenating

df1 = df[['Item_Identifier', 'Item_Weight', 'Item_MRP']].head(10)

df2 = df[['Item_Identifier', 'Outlet_Identifier', 'Item_Outlet_Sales']].head(10)
```

```
# --- Merge (on common column Item_Identifier)

merged = pd.merge(df1, df2, on="Item_Identifier", how="inner")

print("\n==== Merged DataFrame ====")

print(merged)

# --- Join (need indexes)

df3 = df1.set_index("Item_Identifier")

df4 = df2.set_index("Item_Identifier")

joined = df3.join(df4, how="inner")

print("\n==== Joined DataFrame ====")

print(joined)

# --- Concatenate (stacking vertically)

concat_df = pd.concat([df1, df2], axis=0)

print("\n==== Concatenated DataFrame (vertical) ====")

print(concat_df.head(15))
```

Output:-

```

        Outlet_Type  Item_Outlet_Sales
0  Supermarket Type1      3735.1380
1  Supermarket Type2      443.4228
2  Supermarket Type1      2097.2700
3      Grocery Store      732.3800
4  Supermarket Type1      994.7052

==== Column Access ====
0    FDA15
1    DRC01
2    FDN15
3    FDX07
4    NCD19
Name: Item_Identifier, dtype: object

==== Row Access with iloc (first 5 rows, first 3 cols) ====
   Item_Identifier  Item_Weight Item_Fat_Content
0            FDA15       9.30        Low Fat
1            DRC01       5.92        Regular
2            FDN15      17.50        Low Fat
3            FDX07      19.20        Regular
4            NCD19       8.93        Low Fat

==== Row Access with loc (rows where Outlet_Type = 'Supermarket Type1') ====
   Item_Identifier  Item_Weight Item_Fat_Content  Item_Visibility \
0            FDA15       9.30        Low Fat      0.016047
2            FDN15      17.50        Low Fat      0.016760
4            NCD19       8.93        Low Fat      0.000000
6            FD010      13.65        Regular     0.012741
8            FDH17      16.20        Regular     0.016687

   Item_Type  Item_MRP Outlet_Identifier  Outlet_Establishment_Year \
0      Dairy  249.8092          OUT049                  1999
2      Meat  141.6180          OUT049                  1999
4  Household  53.8614          OUT013                  1987
6  Snack Foods  57.6588          OUT013                  1987
8  Frozen Foods  96.9726          OUT045                  2002

   Outlet_Size Outlet_Location_Type      Outlet_Type  Item_Outlet_Sales
0      Medium             Tier 1 Supermarket Type1      3735.1380
2      Medium             Tier 1 Supermarket Type1      2097.2700
4      High              Tier 3 Supermarket Type1      994.7052
6      High              Tier 3 Supermarket Type1      343.5528
8       NaN              Tier 2 Supermarket Type1      1076.5986

==== Null Values Count ====
Item_Identifier           0
Item_Weight                1463
Item_Fat_Content            0
Item_Visibility              0
Item_Type                  0
Item_MRP                    0
Outlet_Identifier            0
Outlet_Establishment_Year      0
Outlet_Size                  2410

```

```

Outlet_Location_Type      0
Outlet_Type                0
Item_Outlet_Sales          0
dtype: int64

==== After Filling Null Values ====
Item_Identifier            0
Item_Weight                 0
Item_Fat_Content            0
Item_Visibility              0
Item_Type                  0
Item_MRP                    0
Outlet_Identifier            0
Outlet_Establishment_Year    0
Outlet_Size                  0
Outlet_Location_Type          0
Outlet_Type                  0
Item_Outlet_Sales            0
dtype: int64

==== Merged DataFrame ====
   Item_Identifier  Item_Weight  Item_MRP Outlet_Identifier  Item_Outlet_Sales
0        FDA15     9.300000  249.8092        OUT049       3735.1380
1        DRC01     5.920000  48.2692        OUT018       443.4228
2        FDN15    17.500000 141.6180        OUT049      2097.2700
3        FDX07    19.200000 182.0950        OUT010       732.3800
4        NCD19     8.930000  53.8614        OUT013       994.7052
5        FDP36    10.395000  51.4008        OUT018       556.6088
6        FDO10    13.650000  57.6588        OUT013      343.5528
7        FDP10    12.857645 107.7622        OUT027      4022.7636
8        FDH17    16.200000  96.9726        OUT045      1076.5986
9        FDU28    19.200000 187.8214        OUT017      4710.5350

==== Joined DataFrame ====
   Item_Identifier  Item_Weight  Item_MRP Outlet_Identifier  Item_Outlet_Sales
Item_Identifier
FDA15             9.300000  249.8092        OUT049       3735.1380
DRC01             5.920000  48.2692        OUT018       443.4228
FDN15            17.500000 141.6180        OUT049      2097.2700
FDX07            19.200000 182.0950        OUT010       732.3800
NCD19            8.930000  53.8614        OUT013       994.7052
FDP36            10.395000  51.4008        OUT018       556.6088
FDO10            13.650000  57.6588        OUT013      343.5528
FDP10            12.857645 107.7622        OUT027      4022.7636
FDH17            16.200000  96.9726        OUT045      1076.5986
FDU28            19.200000 187.8214        OUT017      4710.5350

==== Concatenated DataFrame (vertical) ====
   Item_Identifier  Item_Weight  Item_MRP Outlet_Identifier  Item_Outlet_Sales
0        FDA15     9.300000  249.8092           NaN         NaN
1        DRC01     5.920000  48.2692           NaN         NaN
2        FDN15    17.500000 141.6180           NaN         NaN
3        FDX07    19.200000 182.0950           NaN         NaN
4        NCD19     8.930000  53.8614           NaN         NaN
5        FDP36    10.395000  51.4008           NaN         NaN

```

Practical No 8

```
import pandas as pd
import numpy as np
from sklearn.preprocessing import StandardScaler, MinMaxScaler, LabelEncoder, KBinsDiscretizer
from sklearn.impute import SimpleImputer
# 1. Load Dataset
df = pd.read_csv("Social_Network_Ads.csv")
print("== First 5 rows ==")
print(df.head())
# 2. Handle Missing Values (Imputation)
print("\n== Missing Values Count ==")
print(df.isnull().sum())
# Impute numeric with mean
num_imputer = SimpleImputer(strategy="mean")
df[['Age', 'EstimatedSalary']] = num_imputer.fit_transform(df[['Age', 'EstimatedSalary']])
# Impute categorical with mode
cat_imputer = SimpleImputer(strategy="most_frequent")
df[['Gender']] = cat_imputer.fit_transform(df[['Gender']])
print("\n== After Imputation ==")
print(df.isnull().sum())
# 3. Encoding (Convert Gender to numeric)
le = LabelEncoder()
df['Gender_encoded'] = le.fit_transform(df['Gender'])
print("\n== Encoding Gender ==")
print(df[['Gender','Gender_encoded']].head())
# 4. Standardization (Z-score scaling)
scaler = StandardScaler()
df[['Age_std', 'Salary_std']] = scaler.fit_transform(df[['Age','EstimatedSalary']])
```

```

print("\n==== After Standardization ====")

print(df[['Age','Age_std','EstimatedSalary','Salary_std']].head())

# 5. Normalization (Min-Max scaling)

minmax = MinMaxScaler()

df[['Age_norm','Salary_norm']] = minmax.fit_transform(df[['Age','EstimatedSalary']])

print("\n==== After Normalization ====")

print(df[['Age','Age_norm','EstimatedSalary','Salary_norm']].head())

# 6. Discretization (Binning Age into categories)

discretizer = KBinsDiscretizer(n_bins=4, encode='ordinal', strategy='uniform')

df['Age_bin'] = discretizer.fit_transform(df[['Age']])

print("\n==== After Discretization (Age Bins) ====")

print(df[['Age','Age_bin']].head(10))

# Final Output

print("\n==== Final Processed Data (first 10 rows) ====")

print(df.head(10))

```

Output:-

```

==== First 5 rows ===
   User ID  Gender  Age  EstimatedSalary  Purchased
0  15624510    Male   19          19000        0
1  15810944    Male   35          20000        0
2  15668575  Female   26          43000        0
3  15603246  Female   27          57000        0
4  15804002    Male   19          76000        0

==== Missing Values Count ===
User ID          0
Gender          0
Age            0
EstimatedSalary  0
Purchased       0
dtype: int64

==== After Imputation ===
User ID          0
Gender          0
Age            0
EstimatedSalary  0
Purchased       0

```

```

dtype: int64

==== Encoding Gender ====
   Gender  Gender_encoded
0      Male              1
1      Male              1
2    Female             0
3    Female             0
4      Male              1

==== After Standardization ====
   Age  Age_std  EstimatedSalary  Salary_std
0  19.0   -1.781797     19000.0   -1.490046
1  35.0   -0.253587     20000.0   -1.460681
2  26.0   -1.113206     43000.0   -0.785290
3  27.0   -1.017692     57000.0   -0.374182
4  19.0   -1.781797     76000.0    0.183751

==== After Normalization ====
   Age  Age_norm  EstimatedSalary  Salary_norm
0  19.0    0.023810     19000.0    0.029630
1  35.0    0.404762     20000.0    0.037037
2  26.0    0.190476     43000.0    0.207407
3  27.0    0.214286     57000.0    0.311111
4  19.0    0.023810     76000.0    0.451852

==== After Discretization (Age Bins) ====
   Age  Age_bin
0  19.0      0.0
1  35.0      1.0
2  26.0      0.0
3  27.0      0.0
4  19.0      0.0
5  27.0      0.0
6  27.0      0.0
7  32.0      1.0
8  25.0      0.0
9  35.0      1.0

==== Final Processed Data (first 10 rows) ====

   User ID  Gender  Age  EstimatedSalary  Purchased  Gender_encoded \
0  15624510    Male  19.0     19000.0        0            1
1  15810944    Male  35.0     20000.0        0            1
2  15668575  Female  26.0     43000.0        0            0
3  15603246  Female  27.0     57000.0        0            0
4  15804002    Male  19.0     76000.0        0            1
5  15728773    Male  27.0     58000.0        0            1
6  15598044  Female  27.0     84000.0        0            0
7  15694829  Female  32.0    150000.0        1            0
8  15600575    Male  25.0     33000.0        0            1
9  15727311  Female  35.0     65000.0        0            0

```

	Age_std	Salary_std	Age_norm	Salary_norm	Age_bin
0	-1.781797	-1.490046	0.023810	0.029630	0.0
1	-0.253587	-1.460681	0.404762	0.037037	1.0
2	-1.113206	-0.785290	0.190476	0.207407	0.0
3	-1.017692	-0.374182	0.214286	0.311111	0.0
4	-1.781797	0.183751	0.023810	0.451852	0.0
5	-1.017692	-0.344817	0.214286	0.318519	0.0
6	-1.017692	0.418669	0.214286	0.511111	0.0
7	-0.540127	2.356750	0.333333	1.000000	1.0
8	-1.208719	-1.078938	0.166667	0.133333	0.0
9	-0.253587	-0.139263	0.404762	0.370370	1.0

Practical No 9

```
import pandas as pd
import numpy as np
from sklearn.feature_selection import VarianceThreshold, SelectKBest, f_regression, RFE
from sklearn.decomposition import PCA
from sklearn.linear_model import LinearRegression
# 1. Load Dataset
df = pd.read_csv("auto-mpg.csv")
print("== First 5 rows ==")
print(df.head())
# 2. Handle Missing Values
df.replace('?', np.nan, inplace=True)
# Convert numeric columns to proper dtype
numeric_cols = ['mpg', 'cylinders', 'displacement', 'horsepower', 'weight', 'acceleration', 'model year', 'origin']
for col in numeric_cols:
    df[col] = pd.to_numeric(df[col], errors='coerce')
# Fill missing numeric values with mean
df[numeric_cols] = df[numeric_cols].fillna(df[numeric_cols].mean())
print("\n== Missing values after imputation ==")
print(df.isnull().sum())
# 3. Separate Features and Target
X = df.drop(['mpg', 'car name'], axis=1)
y = df['mpg']
# Ensure all features are numeric
X = X.apply(pd.to_numeric)
# 4. Variance Threshold
var_thresh = VarianceThreshold(threshold=0.1)
X_var = var_thresh.fit_transform(X)
```

```

print("\n==== Variance Threshold Reduction ===")
print("Original features:", X.shape[1])
print("Reduced features:", X_var.shape[1])
# 5. Univariate Feature Selection (SelectKBest)
selector = SelectKBest(score_func=f_regression, k=5)
X_kbest = selector.fit_transform(X, y)
selected_cols = X.columns[selector.get_support()]
print("\n==== Univariate Feature Selection (Top 5) ===")
print(selected_cols)

# 6. Recursive Feature Elimination (RFE)
model = LinearRegression()
rfe = RFE(model, n_features_to_select=5)
X_rfe = rfe.fit_transform(X, y)
rfe_cols = X.columns[rfe.support_]
print("\n==== RFE Selected Features ===")
print(rfe_cols)

# 7. PCA (Principal Component Analysis)
pca = PCA(n_components=2)
X_pca = pca.fit_transform(X)
print("\n==== PCA Components (first 5 rows) ===")
print(X_pca[:5])

# 8. Correlation Analysis
# Only numeric columns for correlation
corr_matrix = df[numERIC_COLS].corr()
print("\n==== Correlation Matrix ===")
print(corr_matrix['mpg'].sort_values(ascending=False))

```

Output:-

```

==== First 5 rows ====
      mpg  cylinders  displacement horsepower  weight  acceleration  model year \
0  18.0          8        307.0       130    3504       12.0         70
1  15.0          8        350.0       165    3693       11.5         70
2  18.0          8        318.0       150    3436       11.0         70
3  16.0          8        304.0       150    3433       12.0         70
4  17.0          8        302.0       140    3449       10.5         70

      origin           car name
0        1  chevrolet chevelle malibu
1        1            buick skylark 320
2        1  plymouth satellite
3        1            amc rebel sst
4        1            ford torino

==== Missing values after imputation ====
mpg          0
cylinders    0
displacement 0
horsepower   0
weight        0
acceleration 0
model year   0
origin        0
car name     0
dtype: int64

==== Variance Threshold Reduction ====
Original features: 7
Reduced features: 7
==== Univariate Feature Selection (Top 5) ====
Index(['cylinders', 'displacement', 'horsepower', 'weight', 'model year'], dtype='object')
==== RFE Selected Features ====
Index(['cylinders', 'horsepower', 'acceleration', 'model year', 'origin'], dtype='object')

==== PCA Components (first 5 rows) ====
[[543.65893755  51.0019395 ]
 [737.54386405  79.3495761 ]
 [478.18600654  75.61166505]
 [473.6069335  62.69978868]
 [488.87747985  55.94417034]]]

==== Correlation Matrix ====
mpg           1.000000
model year    0.579267
origin        0.563450
acceleration  0.420289
horsepower    -0.771437
cylinders     -0.775396
displacement   -0.804203
weight        -0.831741
Name: mpg, dtype: float64

```